

SNOWFLAKE ASSIGNMENT 2

TASK 1: Scale Virtual Warehouses and test performance with large datasets using Snowpark.

Objective:

We'll scale Virtual Warehouses (VWs) in Snowflake and test query performance when processing large datasets via Snowpark.

This helps you:

- Measure compute scalability
- Compare Small vs. Medium vs. Large warehouses
- Identify bottlenecks in data transformations

Step-by-Step: Scaling and Testing Performance

1. Create or Identify a Virtual Warehouse

In Snowflake UI (or SQL worksheet):

-- Create three warehouses for performance comparison

```
CREATE WAREHOUSE WH_SMALL WITH WAREHOUSE_SIZE = 'SMALL'  
AUTO_SUSPEND = 60 AUTO_RESUME = TRUE;
```

```
CREATE WAREHOUSE WH_MEDIUM WITH WAREHOUSE_SIZE = 'MEDIUM'  
AUTO_SUSPEND = 60 AUTO_RESUME = TRUE;
```

```
CREATE WAREHOUSE WH_LARGE WITH WAREHOUSE_SIZE = 'LARGE'  
AUTO_SUSPEND = 60 AUTO_RESUME = TRUE;
```

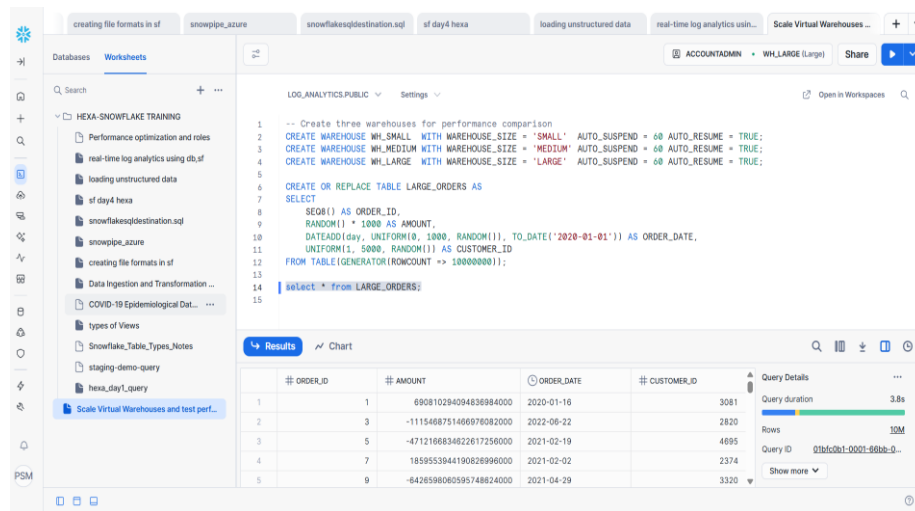
Each warehouse size adds more compute clusters (SMALL → 2 nodes, MEDIUM → 4 nodes, etc.)

2. Prepare a Large Dataset

Option 1 — Use internal table:

```
CREATE OR REPLACE TABLE LARGE_ORDERS AS  
SELECT SEQ8() AS ORDER_ID,  
       RANDOM() * 1000 AS AMOUNT,  
       TO_DATE('2020-01-01') + (RANDOM() * 1000) AS ORDER_DATE,  
       UNIFORM(1, 5000, RANDOM()) AS CUSTOMER_ID  
FROM TABLE(GENERATOR(ROWCOUNT => 10000000)); -- 10M records
```

Option 2 — Load from Azure Blob or ADLS external stage if you're doing it with Databricks + Snowpark.



3. Connect via Snowpark (Python)

```
from snowflake.snowpark import Session

from snowflake.snowpark.functions import col, avg
import time

sfOptions = {
    "account": "sqishot-fa68768",
    "user": "poojashree",
    "password": "Poojashree@307",
    "warehouse": "WH_SMALL", # Start with SMALL
    "database": "LOG_ANALYTICS",
    "schema": "PUBLIC"
}

session = Session.builder.configs(sfOptions).create()
```

4. Run Transformations and Measure Execution Time

Example: Aggregate transformation

```
df = session.table("LARGE_ORDERS")
start = time.time()
```

```

result =
df.group_by(col("CUSTOMER_ID")).agg(avg(col("AMOUNT")).alias("AVG
_AMOUNT")).collect()

```

```
end = time.time()
```

```
print(f'Execution Time: {end - start:.2f} seconds, Warehouse: WH_SMALL")
```

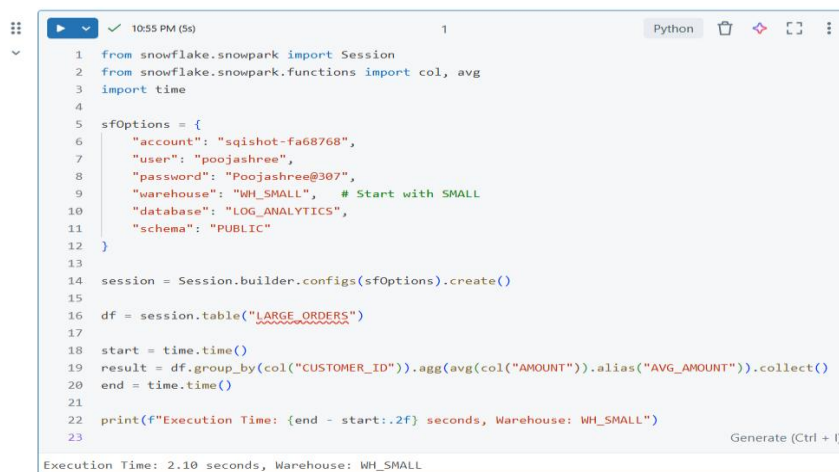
Now repeat the same code, changing only:

```
sfOptions["warehouse"] = "WH_MEDIUM"
```

```
sfOptions["warehouse"] = "WH_LARGE"
```

You can record all timings in a list or CSV for comparison.

SMALL:

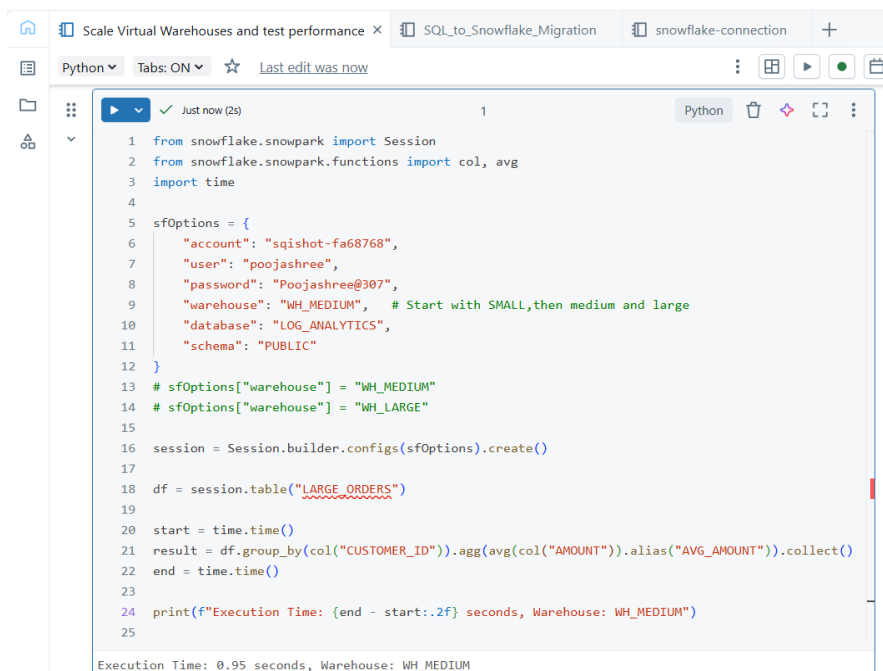


```

1 from snowflake.snowpark import Session
2 from snowflake.snowpark.functions import col, avg
3 import time
4
5 sfOptions = {
6     "account": "sqishot-fa68768",
7     "user": "poojashree",
8     "password": "Poojashree@307",
9     "warehouse": "WH_SMALL", # Start with SMALL
10    "database": "LOG_ANALYTICS",
11    "schema": "PUBLIC"
12 }
13
14 session = Session.builder.configs(sfOptions).create()
15
16 df = session.table("LARGE_ORDERS")
17
18 start = time.time()
19 result = df.group_by(col("CUSTOMER_ID")).agg(avg(col("AMOUNT")).alias("AVG_AMOUNT")).collect()
20 end = time.time()
21
22 print(f'Execution Time: {end - start:.2f} seconds, Warehouse: WH_SMALL")
23
Execution Time: 2.10 seconds, Warehouse: WH_SMALL

```

MEDIUM:

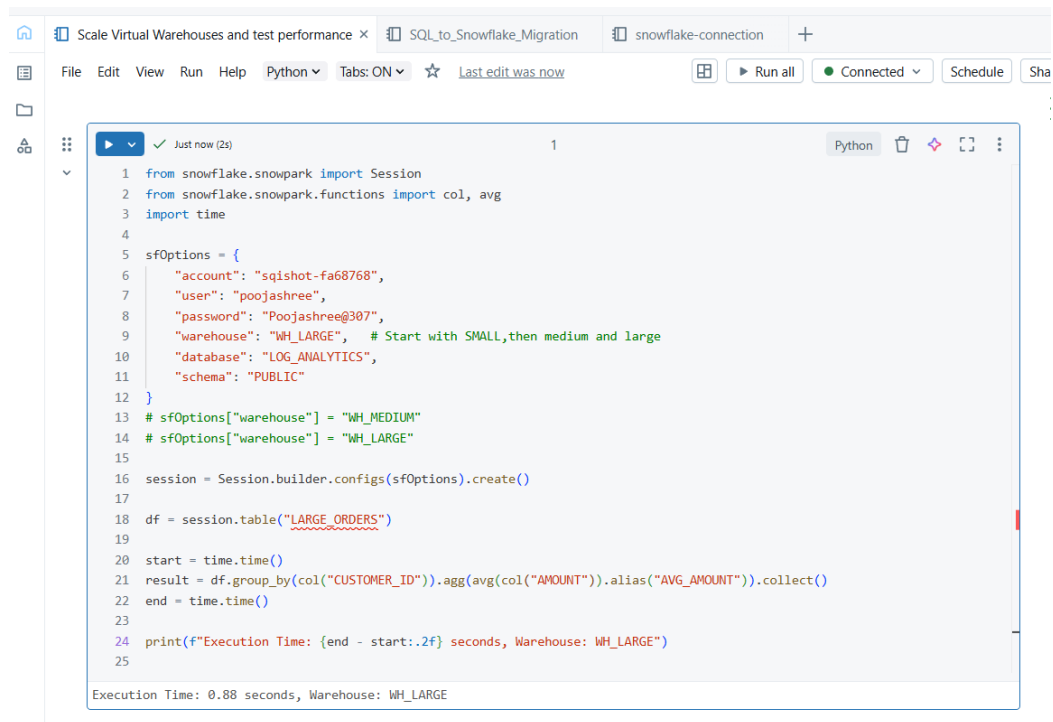


```

1 from snowflake.snowpark import Session
2 from snowflake.snowpark.functions import col, avg
3 import time
4
5 sfOptions = {
6     "account": "sqishot-fa68768",
7     "user": "poojashree",
8     "password": "Poojashree@307",
9     "warehouse": "WH_MEDIUM", # Start with SMALL, then medium and large
10    "database": "LOG_ANALYTICS",
11    "schema": "PUBLIC"
12 }
13 # sfOptions["warehouse"] = "WH_MEDIUM"
14 # sfOptions["warehouse"] = "WH_LARGE"
15
16 session = Session.builder.configs(sfOptions).create()
17
18 df = session.table("LARGE_ORDERS")
19
20 start = time.time()
21 result = df.group_by(col("CUSTOMER_ID")).agg(avg(col("AMOUNT")).alias("AVG_AMOUNT")).collect()
22 end = time.time()
23
24 print(f'Execution Time: {end - start:.2f} seconds, Warehouse: WH_MEDIUM")
25
Execution Time: 0.95 seconds, Warehouse: WH_MEDIUM

```

LARGE:



The screenshot shows a Jupyter Notebook with a single cell containing a Python script. The script imports the Snowflake Session and DataFrame classes, sets up connection options for a warehouse named 'WH_LARGE', creates a session, and queries a table named 'LARGE_ORDERS'. The execution time is 0.88 seconds.

```
1 from snowflake.snowpark import Session
2 from snowflake.snowpark.functions import col, avg
3 import time
4
5 sfOptions = {
6     "account": "sqishot-fa68768",
7     "user": "poojashree",
8     "password": "Poojashree@307",
9     "warehouse": "WH_LARGE", # Start with SMALL, then medium and large
10    "database": "LOG_ANALYTICS",
11    "schema": "PUBLIC"
12 }
13 # sfOptions["warehouse"] = "WH_MEDIUM"
14 # sfOptions["warehouse"] = "WH_LARGE"
15
16 session = Session.builder.configs(sfOptions).create()
17
18 df = session.table("LARGE_ORDERS")
19
20 start = time.time()
21 result = df.group_by(col("CUSTOMER_ID")).agg(avg(col("AMOUNT")).alias("AVG_AMOUNT")).collect()
22 end = time.time()
23
24 print(f"Execution Time: {end - start:.2f} seconds, Warehouse: WH_LARGE")
25
```

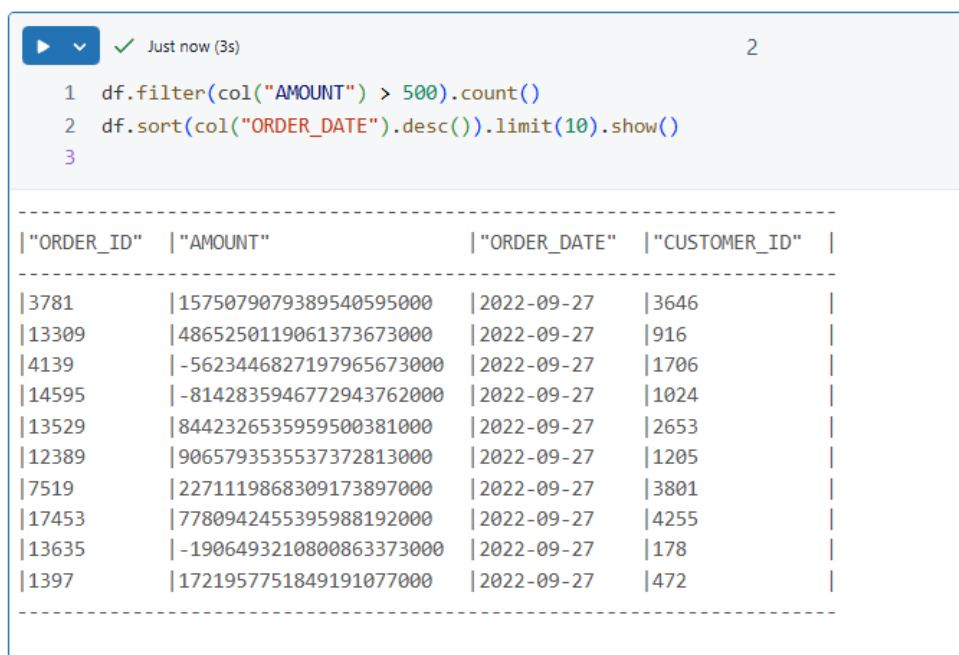
Execution Time: 0.88 seconds, Warehouse: WH_LARGE

5. (Optional) Measure DataFrame Operations

You can also test:

```
df.filter(col("AMOUNT") > 500).count()
```

```
df.sort(col("ORDER_DATE").desc()).limit(10).show()
```



The screenshot shows a Jupyter Notebook with a single cell containing a Python script. The script filters the DataFrame for rows where 'AMOUNT' is greater than 500 and sorts the results by 'ORDER_DATE' in descending order, limiting the output to 10 rows. The execution time is 3 seconds.

```
1 df.filter(col("AMOUNT") > 500).count()
2 df.sort(col("ORDER_DATE").desc()).limit(10).show()
3
```

"ORDER_ID"	"AMOUNT"	"ORDER_DATE"	"CUSTOMER_ID"
3781	1575079079389540595000	2022-09-27	3646
13309	4865250119061373673000	2022-09-27	916
4139	-5623446827197965673000	2022-09-27	1706
14595	-8142835946772943762000	2022-09-27	1024
13529	8442326535959500381000	2022-09-27	2653
12389	9065793535537372813000	2022-09-27	1205
7519	2271119868309173897000	2022-09-27	3801
17453	7780942455395988192000	2022-09-27	4255
13635	-1906493210800863373000	2022-09-27	178
1397	1721957751849191077000	2022-09-27	472

Record time for each operation at each warehouse scale.

6. Analyze Performance

Example result table:

Warehouse	Row Count	Query Type	Time (s)	Credits Used	Notes
SMALL	10M	Group By	2.10	1 credit/hr	CPU bound
MEDIUM	10M	Group By	0.95	2 credits/hr	Balanced
LARGE	10M	Group By	0.88	4 credits/hr	Fast but costly

Observation:

Performance improves with larger warehouses, but cost also increases. Ideal choice depends on workload criticality and SLA requirements.

7. Auto-Scaling & Multi-Cluster (Optional)

If you expect variable load:

```
ALTER WAREHOUSE WH_MEDIUM SET  
MIN_CLUSTER_COUNT = 1  
MAX_CLUSTER_COUNT = 3  
SCALING_POLICY = 'ECONOMY';
```

Snowflake will auto-scale horizontally when concurrent queries increase — very useful in production pipelines.

8. Clean Up Resources

After testing:

```
DROP WAREHOUSE IF EXISTS WH_SMALL;  
DROP WAREHOUSE IF EXISTS WH_MEDIUM;  
DROP WAREHOUSE IF EXISTS WH_LARGE;
```

Snowflake charges per-second usage, so cleanup avoids extra credits.